

AML PROGRAM UPDATE — SECTION 4.c

Investigation Metrics & Quality Assurance Framework

SLA Tiers · QA Sampling Methodology · Rule Precision & Recall Measurement

Prepared by: AML Program Manager | Scope: Transaction Monitoring & Investigations |
Reference dataset period: 2025-03-01 to 2025-11-07 (252 days)

*Candidate: Vitor Cruz | Submission date: 2026-08-24 | Target role: AML/CTF Analyst (US) —
CloudWalk Payments Inc.*

Executive Summary

This document establishes the operating framework that governs how CloudWalk's AML Transaction Monitoring and Investigations function triages, works, and quality-controls alerts, and how the underlying detection rules are measured and recalibrated over time. It is built directly on the transaction population confirmed in the case dataset (32,348 transactions across 252 days; 4,078 declines; 236 chargebacks; 350 refunds; 41 high-risk merchants; 16 PEP customers; 1 customer above the 0.7 sanctions-match threshold) and on the alert typologies and thresholds already defined in the Section 4.a detection queries and the Section 1.b program update.

Four components are defined below: (1) a three-tier SLA structure (P1/P2/P3) that classifies alerts by severity and sets time-bound targets for acknowledgment, investigation, and disposition; (2) a QA sampling methodology that scales review intensity to tier risk and evaluates both case quality and investigator consistency; (3) a precision/recall measurement framework that turns chargeback outcomes, SAR filings, and look-back testing into a repeatable monthly rule-tuning cycle; and (4) an investigation capacity model that converts the tier structure into a staffing requirement and ties a phased false-positive-rate reduction target to the rule corrections already delivered in this case file. All four feed the quarterly CloudWalk-Lead Bank governance committee established under Section 9 of the AML program.

1. Investigation SLA Tiers (P1 / P2 / P3)

Tier assignment is automated at alert generation, driven by the same fields used in the Section 4.a detection queries: `pep_flag`, `sanctions_match_score`, `high_risk_mcc_flag`, `amount_usd` (relative to `expected_avg_ticket_usd`), and the device-sharing / chargeback-ratio thresholds already codified in the program. An alert is assigned the highest tier triggered by any single criterion — tiers are not averaged.

1.1 Tier Summary Table

Tier	Classification Criteria (any one triggers the tier)	Alert Acknowledgment SLA	Full Investigation SLA	Decision SLA (close / escalate / SAR)
P1 — Critical	- sanctions_match_score > 0.70 - pep_flag = TRUE AND high_risk_mcc_flag = TRUE (combined PEP + high-risk-MCC hit) - VE or IR nexus via country or doc_issue_country (mandatory KYC upgrade trigger) - Device-sharing hard-block: ≥5 distinct customer_ids on one device_id within 60 minutes - Merchant chargeback ratio ≥ 8% (30-day, auto-suspend tier)	1 hour	24 hours	48 hours (aligns with the 48-hour Lead Bank BSA Officer escalation SLA, Section 9)
P2 — High	- pep_flag = TRUE (standalone, not already captured in P1) - high_risk_mcc_flag = TRUE AND amount_usd > 2x expected_avg_ticket_usd - Merchant chargeback ratio	4 business hours	5 business days	7 calendar days

Tier	Classification Criteria (any one triggers the tier)	Alert Acknowledgment SLA	Full Investigation SLA	Decision SLA (close / escalate / SAR)
	3%–5% (yellow) or >5%–8% (red) over rolling 30 days • Device-sharing alert: 3–4 distinct customer_ids on one device_id within 60 minutes • IP ring: ≥3 distinct customer_ids on one shared ip_address, ECOM channel			
P3 — Standard	• Single-signal, lower-severity hits not otherwise escalated to P1/P2: • ECOM + KEYED (no 3DS) at a high_risk_mcc_flag = TRUE merchant • Structuring-band transaction at a high-risk MCC with no PEP/sanctions overlay • Low-value card-testing pattern (ECOM, amount_usd < \$5.00) • Routine high-risk-MCC transaction with no other risk	1 business day	10 business days	15 calendar days

Tier	Classification Criteria (any one triggers the tier)	Alert Acknowledgm ent SLA	Full Investigation SLA	Decision SLA (close / escalate / SAR)
	flag present			

1.2 Worked Examples from the Dataset

P1 example — the confirmed Iranian-passport customer with `sanctions_match_score` = 0.84 (referenced in the Section 4 typology analysis) triggers P1 on the sanctions criterion alone; the confirmed 12-distinct-customer device-sharing incident at MCC 4829 independently triggers P1 on the hard-block criterion, exceeding the ≥ 5 -customer threshold by more than double.

P2 example — a PEP customer transacting at a `high_risk_mcc_flag` = TRUE merchant (MCC 4829) at 3.6x their `expected_avg_ticket_usd` (illustrated in the Section 4.a PEP/high-risk-MCC query, e.g., customer C10099) triggers P2 on the PEP + ticket-multiple criterion. A merchant sitting at a 4% chargeback ratio (above the 3% yellow-alert line but below the 5% red-escalation line) also enters at P2.

P3 example — an otherwise unremarkable customer with no PEP, sanctions, or device/IP flag, transacting ECOM with `pos_entry_mode` = KEYED at a high-risk-MCC merchant, is routed to P3 for routine EDD-adjacent review rather than urgent escalation. The confirmed outlier merchant M3030 (ElectroHub #330) at a 12% chargeback ratio does not appear at P3 — it clears the P1 auto-suspend threshold and is handled as Critical.

1.3 Operating Notes

- SLA clocks start at alert generation timestamp, not at analyst login; business-hour tiers (P2/P3) run on the investigations desk's published business calendar, while P1's hour-based clock runs 24/7 given its severity.
- A P3 alert that develops a second, higher-severity signal during investigation (e.g., a KEYED-ECOM case where the customer is subsequently confirmed PEP) is immediately re-tiered to P2/P1 and its SLA clock restarts from the re-tier timestamp, not the original alert time.
- SLA breaches are logged automatically in the case management system and roll up into the monthly QA report (Section 2) and the quarterly governance committee packet (Section 9 of the AML program).

2. QA Sampling Methodology

QA sample intensity is inversely scaled to alert volume and directly scaled to severity: the rarer, higher-stakes P1 population receives full review, while the higher-volume P3 population is reviewed on a statistically representative sample supplemented by targeted pulls.

2.1 Sampling Frequency and Selection by Tier

Tier	Sampling Frequency	Sample Basis	Selection Criteria
P1	Continuous — every closed case	100% of population	Full review; no sampling applied given low volume and severity. Every P1 disposition is reviewed by the QA lead within 5 business days of closure, prior to being reported to the quarterly governance committee.
P2	Monthly	~20% of closed cases (10% random + 10% targeted)	Random: stratified random draw across analysts and typologies so no single analyst or rule dominates the sample. Targeted: all cases closed as “no further action” that involved a PEP or a merchant above the 3% chargeback threshold; all cases where the investigator’s decision was overturned on a prior QA review; all cases tied to the confirmed outlier merchant (M3030) or to a repeat device/IP identifier.
P3	Monthly	~5% of closed cases (random)	Simple random sample stratified by typology (structuring-band, ECOM-no-3DS, card-testing) so each typology is

Tier	Sampling Frequency	Sample Basis	Selection Criteria
			represented roughly in proportion to its alert volume; a minimum of 3 cases per active P3 typology per month regardless of volume, to avoid a low-volume typology going unreviewed.

New analysts (first 90 days) and any analyst who fell below the 90% concordance target in the prior month (Section 2.3) are sampled at 100% regardless of tier, until two consecutive clean monthly reviews are achieved.

2.2 What Is Evaluated in Each Sampled Case

- SLA adherence — acknowledgment, investigation, and decision timestamps against the applicable tier target, with reason codes captured for any breach.
- Correct tier classification — whether the alert was assigned and, if applicable, re-tiered in line with Section 1's criteria.
- Evidence sufficiency — whether the case file supports the disposition with transaction-level detail (txn_id, amount_usd, mcc, counterparties) rather than conclusory statements.
- Narrative quality — for cases with a written narrative, structure and completeness against the case playbook's Who/What/When/Where/How/Counterparties/Fund flows/Prior history/Recommendation/Rationale format.
- Decision appropriateness — whether close / escalate / SAR is consistent with the facts, the typology definitions, and the program's mandatory-escalation triggers (e.g., sanctions_match_score > 0.7 must not be closed without EDD/escalation).
- Recordkeeping completeness — presence of the mandatory fields defined in Section 6 of the program update (transaction, CIP, and KYB fields) in the retained case file.
- Consistency across analysts — whether two analysts presented with materially similar fact patterns reached materially similar dispositions.

2.3 Quality Metrics

Metric	Definition	Target
QA pass rate	Share of sampled cases with no material finding (no SLA breach, correct tier, correct	≥ 95%

Metric	Definition	Target
Investigator-QA concordance rate	decision, adequate evidence). Share of sampled cases where the QA reviewer independently agrees with the investigator's close/escalate/SAR decision.	≥ 90%
Narrative quality score	1–5 rubric score (5-Ws-and-H completeness, evidence citation, clarity of SAR rationale) averaged across sampled narratives.	≥ 4.0 / 5.0
Rework rate	Share of sampled cases returned to the investigator for revision before the QA finding is finalized.	≤ 8%
SLA timeliness compliance	Share of closed cases, by tier, meeting the acknowledgment, investigation, and decision SLAs in Section 1.	≥ 95% P1 / ≥ 90% P2 / ≥ 85% P3

QA findings below target trigger root-cause review (training gap vs. rule/threshold ambiguity vs. workload) and feed both the analyst coaching cycle and the rule-tuning cycle described in Section 3.

3. Precision and Recall Measurement for Detection Rules

3.1 Definitions in This Context

True Positive (TP) — an alert whose investigation confirms genuinely suspicious activity, evidenced by any of: escalation to the Lead Bank BSA Officer, a SAR filed, a documented risk-rating upgrade or account/merchant restriction, or closure with a formal EDD referral. A TP does not require a SAR to have been filed — restrictive action taken on the strength of the alert also counts, since not every confirmed risk finding meets the SAR filing threshold.

False Positive (FP) — an alert closed as “no suspicious activity identified” after investigation produces a documented, verifiable legitimate explanation (e.g., confirmed business purpose, KYC-consistent activity), with no repeat trigger on the same customer/merchant within a 90-day look-back window. A repeat trigger within 90 days reclassifies the original disposition as a miss for root-cause review, not a clean FP.

False Negative (FN) — suspicious activity later confirmed through a channel other than the rule in question — a chargeback pattern, a 314(a) law-enforcement name match, an external tip, a SAR filed on related activity, or QA look-back sampling of the non-alerted population —

where the underlying transactions were present in the dataset and met conditions that, in hindsight, should have triggered an alert but did not.

Because the universe of undetected suspicious activity is not directly observable, recall is measured as a proxy rather than an exact rate (Section 3.3).

3.2 Tuning Frequency

- Formal quarterly recalibration of every rule's thresholds, timed to the quarterly CloudWalk–Lead Bank governance committee (Section 9 of the AML program), using the trailing quarter's precision/recall data.
- Ad hoc tuning within 30 days of detection whenever a rule's trailing-90-day precision falls below its floor (10% for single-signal behavioral rules such as ECOM-no-3DS or structuring-band; 5% for combination/network rules such as PEP + high-risk MCC or device/IP rings, which are intentionally lower-precision, higher-recall rules) or whenever a chargeback- or SAR-driven false negative reveals a materially under-calibrated threshold.
- Every threshold change is logged in a change-management record: prior value, new value, supporting data, and CCO approval, consistent with the parameter-reference documentation already maintained alongside the Section 4.a detection queries.

3.3 Using Chargeback and SAR Data to Calibrate

Chargebacks as a false-negative signal. Any confirmed chargeback (`is_chargeback = TRUE`) on a customer or merchant that had no prior alert is treated as a candidate false negative and is fed into the monthly review. Because the dataset's baseline chargeback ratio is ~0.73% (236 of 32,348 transactions) against a confirmed 12% outlier at merchant M3030, chargeback trend is also read as a leading-indicator test: if a merchant's chargeback ratio crosses the 3% yellow-alert threshold before any transaction monitoring alert fired on that merchant, the monitoring rule is treated as lagging and its trigger conditions are tightened.

SAR filings as a true-positive and yield signal. Every filed SAR is tagged back to the rule(s) or typology(ies) that produced the originating alert(s). Where a SAR draws on more than one typology (e.g., a case combining PEP-at-high-risk-MCC and device-sharing evidence), each contributing rule receives credit toward its TP count, consistent with the case playbook's practice of documenting every typology present in an alert rather than a single primary trigger. This produces a per-rule "SAR yield" — SARs attributable to the rule divided by alerts the rule generated — as well as a program-level SAR-to-alert ratio tracked over time.

Recall proxy construction. True recall is unmeasurable without a labeled ground truth, so it is approximated from four converging signals: (1) chargeback-driven back-testing described above; (2) periodic look-back QA sampling of the non-alerted transaction population, re-run against current thresholds; (3) 314(a)/law-enforcement matches against customers with no prior alert history; and (4) post-SAR-filing review of the filed subject's related, unflagged transactions and counterparties. A rule is flagged for recall concern when any of these four sources repeatedly surfaces activity the rule's current thresholds do not catch.

3.4 Monthly Rule Effectiveness Report — Format

Published monthly to the AML Program Manager and quarterly (rolled up) to the CloudWalk–Lead Bank governance committee. One row per rule/typology:

Rule / Typology	Alerts Generated				Chargeback - Linked FNs Found		Recal Proxy Signal	SARs Attribute	SAR Yield %	Tuning Action	Owner
		TP	FP	Precision %							
e.g., PEP + High-Risk MCC	—	—	—	—	—	—	—	—	—	—	—

Template rows are populated monthly for each Section 4.a rule (chargeback-ratio tiering, PEP/high-risk-MCC, device sharing, IP ring, ECOM-without-3DS, structuring-band, FATF/OFAC jurisdiction, card testing). The report closes with a short narrative covering: the highest-precision and lowest-precision rule of the month, any threshold changes made or proposed, chargeback- or SAR-driven false negatives discovered, and the plan for next month’s recalibration.

5. Investigation Capacity Model

This section translates the P1/P2/P3 tier structure defined in Section 1 into a staffing requirement, using the confirmed alert population from the Section 4.b typology analysis (32,348 transactions / 501 customers / 100 merchants / 630 cards; 252-day period) as the volume basis. It closes the loop between detection design (Section 1), quality control (Section 2), rule tuning (Section 3), and the resourcing needed to run all three on an ongoing basis.

5.1 Data Basis and Methodology

The dataset’s confirmed, typology-tagged alert population totals 676 records over the confirmed 252-day reporting period (2025-03-01 to 2025-11-07 inclusive; 251 elapsed days between first and last transaction) — ≈81.7 alerts/month, per the Section 6.b KPI dashboard baseline, summing the eleven typology-confirmation counts from the Section 4.b analysis (T01–T11). Two of those eleven typologies — T03 (Device Sharing) and T09 (IP Ring) — are separate rule tags on the same underlying incident: the 12-linked-customer mule ring is one investigative case that happens to trigger both the device-sharing hard-block and the IP-ring criterion simultaneously. Typology tracking counts it twice (once per rule) for rule-effectiveness reporting; a case management workload does not, since one investigator works it once regardless of how many rules fired. The capacity model below therefore uses 675 unique investigative cases (676 minus the one linked-case duplicate) as its volume denominator,

consistent with Section 3.3's practice of crediting a single alert to every contributing typology without treating it as multiple units of investigator work.

Each typology is mapped to the tier it triggers under the Section 1.1 criteria table (highest-tier-wins, as specified in Section 1):

Typology (Section 4.b)	Confirmed Count	Tier Triggered	Section 1.1 Criterion Matched
T11 — FATF/OFAC Jurisdiction	3	P1	sanctions_match_score > 0.70 (1 Iranian-passport customer, score 0.84) and VE/IR nexus via doc_issue_country (2 Venezuelan customers, FATF Grey List)
T05 — PEP + High-Risk MCC	16	P1	pep_flag = TRUE AND high_risk_mcc_flag = TRUE (combined hit) — all 16 confirmed PEP customers transact at high-risk-MCC merchants, per the Section 4.b Summary
T07 — Chargeback Outlier	1	P1	Merchant chargeback ratio ≥ 8% (auto-suspend) — M3030 confirmed at 12.0%, the only merchant in the portfolio above the 3% yellow-alert line (direct query against the 100-merchant population found zero merchants in the 3%–8% P2 band)
T03 / T09 — Device-Sharing / IP-Ring incident	1 (linked case)	P1	Device-sharing hard-block, ≥5 distinct customer_ids on one device within 60 minutes — 12 confirmed, more

Typology (Section 4.b)	Confirmed Count	Tier Triggered	Section 1.1 Criterion Matched
T08 — Cash-In/Cash-Out (tightened, ≤72h, ≥80% ratio)	1	P2 (mapped by analogy)	than double the threshold Not separately named in Section 1.1; treated as an elevated single-signal behavioral typology given its high-confidence filter, pending an explicit Section 1.1 line item (see Recommendation in 5.5)
T01 — Structuring	6	P3	Structuring-band transaction at a high-risk MCC with no PEP/sanctions overlay
T02 — Card Testing	8	P3	Low-value card-testing pattern (ECOM, amount_usd < \$5.00)
T04 — Cross-Border / Geo-Hopping	342	P3 (mapped by analogy)	Not separately named in Section 1.1; routed to the P3 catch-all (“routine high-risk-MCC transaction with no other risk flag present”) since cross-border activity alone, absent a VE/IR nexus or a second overlay, does not meet any named P1/P2 criterion
T06 — ECOM without 3DS	297	P3	ECOM + KEYED (no 3DS) at a high_risk_mcc_flag = TRUE merchant

Two verification checks were run directly against the transaction data before finalizing this table, rather than assuming the P2 criteria were populated: a full portfolio scan confirmed no merchant sits in the 3%–8% chargeback band, and a 60-minute sliding-window scan (the same window logic as the Section 4.a Query 2 design) confirmed no device or shared-IP cluster reaches the 3–4-customer P2 band independent of the one confirmed 12-customer P1 ring. Both P2 criteria checks came back empty against the confirmed population; this is a finding in its own right and is discussed in 5.5.

5.2 Monthly Alert Volume by Tier

Tier	Unique Cases (252-day period)	Share of Volume	Estimated Monthly Volume
P1 — Critical	21	3.1%	≈2.54 / month
P2 — High	1	0.1%	≈0.12 / month
P3 — Standard	653	96.7%	≈78.88 / month
Total	675	100%	≈81.54 / month

The split is heavily bimodal: confirmed severe activity clears a named P1 trigger outright (sanctions, PEP-at-high-risk-MCC, the chargeback auto-suspend line, or the device/IP hard-block), while nearly everything else that survived to “confirmed” status in the Section 4.b analysis is a routine, single-signal P3 hit. The near-empty P2 tier is a genuine finding about this rule set and dataset, not a rounding artifact — see 5.5.

5.3 Average Handling Time Assumptions

The figures below are planning estimates drawn from general BSA/AML investigations staffing benchmarks, not empirical time-and-motion measurements of CloudWalk’s own case files; the case management system should begin capturing actual open-to-close analyst hours per tier so these can be replaced with measured figures at the first quarterly recalibration.

Tier	Assumed Avg. Handling Time	Basis
P1	5.0 hours	Full investigation-to-decision cycle for a case meeting the 24-hour investigation SLA (Section 1.1), including OFAC/EDD coordination, device/IP forensic documentation, and a SAR-ready narrative per the case playbook’s Who/What/When/Where/How/Counterparties/Fund-flows/Prior-history/Recommendation/Rationale format (Section 2.2). Set above the user’s initial 4-hour estimate

Tier	Assumed Avg. Handling Time	Basis
		to reflect the documented complexity of the Section 3/4.4 SAR evidence packages already produced in this case file for comparably severe fact patterns; treat as a 4–6 hour range pending measured data.
P2	1.5 hours	Standard investigation within the 5-business-day SLA; single-signal-plus-overlay cases requiring EDD-adjacent review but not full SAR drafting.
P3	0.5 hours	Routine single-signal disposition within the 10-business-day SLA; high-volume, low-complexity review consistent with the tier’s “routine” classification in Section 1.1.

5.4 FTE Requirement Calculation

FTE = [Σ (monthly volume per tier $\times$ handling time per tier)] $\div$ (productive hours/month $\times$ utilization factor)

Monthly investigator-hours required, by tier:

Tier	Monthly Volume	Handling Time	Monthly Hours
P1	2.54	5.0h	12.7
P2	0.12	1.5h	0.2
P3	78.88	0.5h	39.4
Total	81.54	—	52.3 hours/month

Available hours per FTE: a standard 40-hour week yields 173.3 gross monthly hours. Netting out PTO, holidays, training, team meetings, and coaching (~86% retention, an industry-standard planning ratio for specialist compliance roles) leaves an estimated 150 productive hours/month. Applying the requested 0.75–0.80 utilization factor — the share of productive time actually available for direct casework once QA remediation, case-file administration beyond narrative drafting, and system/process overhead are set aside — yields 112.5–120 available investigation hours per FTE per month.

Utilization Factor	Available Hours/FTE/Month	FTE Required
0.75	112.5	0.46
0.775 (midpoint)	116.25	0.45
0.80	120.0	0.44

Computed requirement: ≈ 0.44 – 0.46 FTE dedicated to tiered alert investigation at current confirmed volume.

This is a floor, not a staffing recommendation. A sub-1.0 FTE result does not mean CloudWalk can run this program on a fraction of an analyst: it means alert-investigation load alone does not fill one person’s time. A functioning program still needs:

- **A minimum of 1.0 dedicated Investigations Analyst FTE** for continuity — coverage during PTO, onboarding lead time, and a single point of failure otherwise on an EDD program that already carries 242 open EDD cases (Section 2.a) and a growing transaction base.
- **QA lead time**, separately budgeted per the Section 2.1 sampling table (100% of P1, $\sim 20\%$ of P2, $\sim 5\%$ of P3), which this calculation does not include.
- **Headroom for volume growth** — the 81.54/month baseline reflects the confirmed dataset as measured; live monitoring volume should be expected to grow with card/merchant onboarding and is not a static ceiling.

The remaining $\approx 0.5+$ FTE of unused alert-investigation capacity is the natural home for EDD population remediation (Section 2.a’s 242 cases), QA second-line review, and the rule-tuning cycle in Section 3 — work this program already has queued and unstaffed, rather than idle time.

5.5 Phased False-Positive-Rate Reduction Target

No full-population precision/recall measurement exists yet for the 675-case confirmed population — Section 3’s framework is prospective, and the only case-level dispositions produced to date are the curated Section 4.4 Top-10 (9 SAR recommendations, 1 escalation, effectively 100% TP by construction of that sample). A programwide FPR baseline is therefore not yet a measured number, and the target below is anchored instead to two things that are already known: the Section 3.2 precision floors, and a concrete result already delivered by Task 7.4.

Task 7.4 as the reference case. The T10 (Self-Merchant / undisclosed-UBO) rule, before correction, would have generated an alert for every one of 175 name-only UBO-to-customer matches — and 0 of those 175 are corroborated by a second identity attribute (date of birth), against a ~ 1 -in-40 chance-collision rate expected from the dataset’s shared name pool alone. Left uncorrected, T10 would have run at an effective 0% precision (100% FPR), adding an estimated 21.1 pure-noise alerts/month — a 25.9% inflation over the current 81.54/month baseline, and roughly 10.6 additional investigator-hours/month (≈ 0.09 FTE) for zero investigative yield. Task 7.4’s fix — requiring a second, independent corroborating attribute

before the rule fires, rather than a single weak signal (name match alone) — already removed this entirely; the 81.54/month baseline used throughout this section is the corrected, post-Task-7.4 figure.

Applying the same principle forward. The same corroboration-first pattern is proposed for the four confirmed single-signal P3 rules that make up 96.7% of current volume (T01 Structuring, T02 Card Testing, T04 Cross-Border, T06 ECOM-without-3DS), each currently operating at the Section 3.2 regulatory floor for single-signal rules (10% precision / up to 90% FPR before ad hoc tuning is triggered):

Phase	Timing	Target	Mechanism	Estimated Effect
Phase 1	Next quarterly recalibration (Section 3.2)	Raise trailing-90-day precision on the four P3 single-signal rules from the 10% floor to $\geq 15\%$	Add one corroborating condition to each rule (e.g., cross-border AND a second high-risk-MCC/amount overlay for T04; tighten the KEYED/no-3DS threshold for T06), following the same second-attribute discipline validated by Task 7.4; changes logged per the Section 3.2 change-management record	Directional target of 10–15% fewer P3 alerts/month (≈ 8 –12 fewer alerts), re-validated against the Section 3.3 chargeback/SAR recall-proxy signals before being finalized, so recall is not sacrificed for precision
Phase 2	Second quarterly recalibration (+2 quarters from Phase 1)	Raise the same four rules to $\geq 20\%$ precision	Re-tune based on two full quarters of monthly rule-effectiveness reports (Section 3.4) and confirmed recall-proxy data, rather than a single-	Directional target of 20–25% cumulative reduction in P3 volume from the current baseline (≈ 16 –20 fewer alerts/month), freeing an estimated 8–10 investigator-

Phase	Timing	Target	Mechanism	Estimated Effect
			cycle adjustment	hours/month (≈0.07–0.09 FTE) for EDD and QA work

Both phases are targets to validate against live data at the stated recalibration points, not commitments made in the absence of that data — consistent with Section 3.2’s own ad hoc/quarterly tuning cadence and the precision floors already governing when a rule must be revisited. Each phase’s outcome should be reported in the Section 3.4 monthly format and rolled into the quarterly governance packet (Section 6) alongside the SLA and QA metrics.

Recommendation carried into the next Section 1.1 revision: T04 (Cross-Border) and T08 (Cash-In/Cash-Out) are confirmed, actively-tracked typologies with no explicit line item in the Section 1.1 tier criteria table; both were mapped by analogy for this capacity model (5.1). The next tier-table revision should give them named criteria so tier assignment for these two typologies is rule-driven rather than case-by-case judgment.

6. Governance Integration

SLA compliance, QA findings, the monthly rule-effectiveness report, and the capacity model’s FTE and phased FPR-reduction targets (Section 5) are consolidated into a single quarterly package for the CloudWalk–Lead Bank governance committee alongside filing volumes and the 48-hour escalation SLA already tracked under Section 9 of the AML program. This keeps investigation quality, alert-tier discipline, rule calibration, and staffing adequacy on the same reporting cadence as the program’s existing regulatory governance cycle, rather than as a parallel, disconnected metric set.